



KESHAV MEMORIAL INSTITUTE OF TECHNOLOGY


AN AUTONOMOUS INSTITUTION - ACCREDITED BY NAAC WITH 'A' GRADE

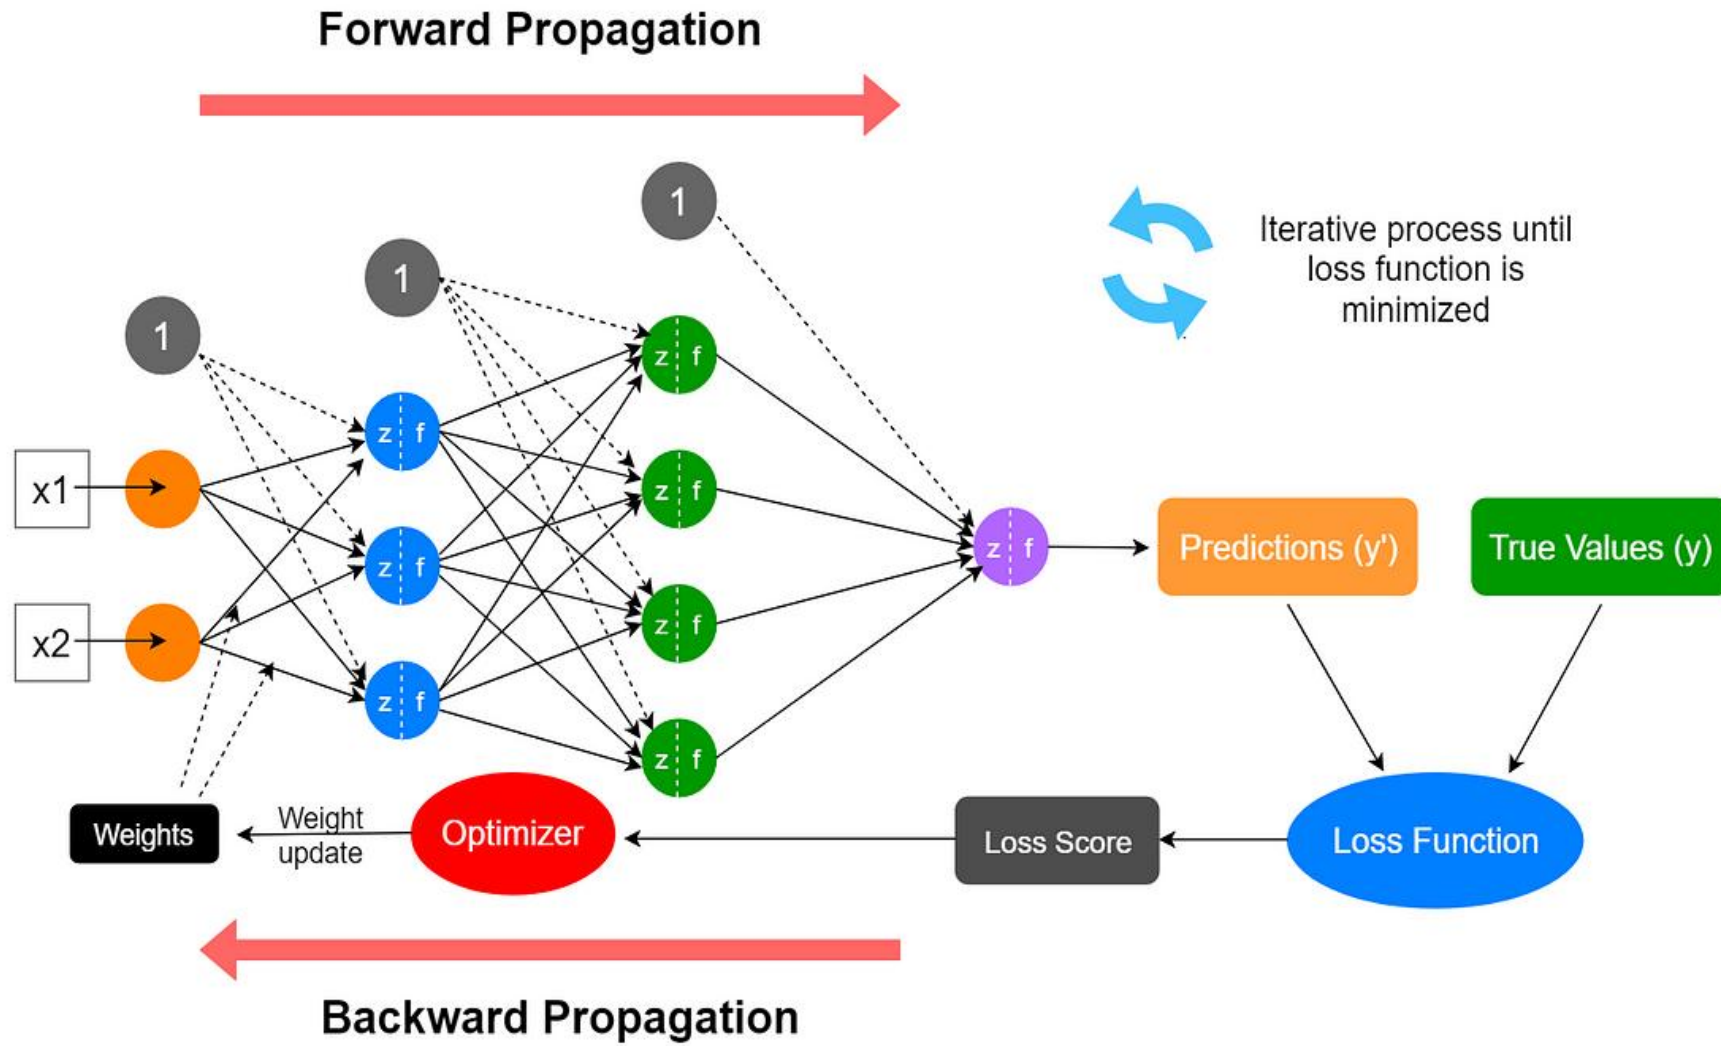
Narayanaguda, Hyderabad.

NATURAL LANGUAGE PROCESSING (21CM601PC)

**BY
ASHA
ASSISTANT PROFESSOR
CSE(AI&ML)
KMIT**

Activation Functions in Neural Networks





Forward Propagation

- The process of computing the output of a neural network. It involves:
 1. Multiplying inputs by weights.
 2. Adding biases.
 3. Applying the activation function to produce the output.

Why do we need Activation functions?

Activation functions are crucial in neural networks because they introduce **non-linearity** into the model, enabling the network to learn and model complex patterns in the data. Without activation functions, a neural network would essentially behave like a linear regression model, limiting its ability to capture the true underlying structure in most real-world problems.

$$Z1 = W1 * A0 + B1$$

~~$$A1 = f(Z1)$$~~

$$Z2 = W2 * A1 + B2$$

~~$$A2 = f(Z2)$$~~

$$Z3 = W3 * A2 + B3$$

~~$$A3 = f(Z3)$$~~

$$A2 = W2 * (W1 * A0 + B1) + B2$$

$$A2 = \underbrace{W2 * W1}_{W'} * A0 + \underbrace{W2 * B1 + B2}_{B'}$$

$$A2 = W' * A0 + B'$$

$$A3 = W'' * A0 + B''$$

1. Real-Life Analogy: The Light Dimmer

- Without activation → neurons behave like a simple *on/off switch* (just linear scaling).
 - With activation → neurons behave like a **dimmer** switch → brightness adjusts differently at different levels.
- 👉 This shows that activation functions give neurons *nonlinear control*, not just a fixed slope.

1. How real neurons work

- A neuron in the brain receives inputs from thousands of other neurons.
 - Each input has a *weight* (strength of connection).
 - The neuron **adds up all these inputs**.
 - If the total **crosses a threshold**, the neuron **fires** (sends a signal forward).
 - Otherwise, it stays silent.
- 👉 So, the brain is **not linear** → it's full of **nonlinear thresholding and modulation**.

2. Artificial Neurons mimic this

In Artificial Neural Networks (ANNs), each neuron also **sums up inputs**:

$$z = w_1 \cdot x_1 + w_2 \cdot x_2 + \dots + b$$

But if we just pass z forward as it is \rightarrow it's linear.

So we need an **activation function** to decide:

- Should the neuron fire strongly?
- Should it suppress weak signals?
- Should it pass values partially (not just ON/OFF)?

“Activation functions in ANNs exist because, just like biological neurons, we don’t want every signal to pass through linearly. We want the neuron to **decide how much to fire** – weakly, strongly, or not at all.”

Non-Linearity

- Why needed? → To learn complex, nonlinear patterns in data

Step Function (Threshold Function)

- Formula:

$$f(x) = 1 \text{ if } x > 0, \text{ else } 0.$$

Intuition: Neuron fires (1) or stays silent (0).

$$f(z) = \begin{cases} 1 & \text{if } z \geq \theta \\ 0 & \text{if } z < \theta \end{cases}$$

- $z=wx+b$ is the neuron's pre-activation (weighted sum + bias).

- **Use:**
 - Early **perceptrons** (1950s–60s).
 - Rarely used today → not differentiable, so **not good for backpropagation.**
- **When?** Only for very simple **binary decision rules.**

ReLU (Rectified Linear Unit)

Formula:

$$f(x) = \max(0, x)$$

Derivative:

$$f'(x) = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases}$$

- **Intuition:** Pass positive signals, block negatives (like a switch).

Use:

Most popular activation for hidden layers.

Simple, efficient, avoids vanishing gradient (mostly).

Limitations:

“Dying ReLU” problem → neurons stuck at 0.

When?

Default choice in **hidden layers of deep networks**.

Dying ReLU problem happens when a neuron **always outputs 0** for all inputs.

This means the neuron is effectively **dead**: it never activates, never contributes to learning.

The gradient through it is also **0** (since slope for $x < 0$ is zero).

Once dead, it **never recovers**, because weight updates depend on gradients.

How to fix dying ReLU?

Leaky ReLU:

Leaky ReLU / Parametric ReLU

Formula:

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha x & \text{if } x \leq 0 \end{cases}$$

Derivative:

$$f'(x) = \begin{cases} 1 & \text{if } x > 0 \\ \alpha & \text{if } x \leq 0 \end{cases}$$

Intuition: Fixes dead neurons problem.


Use:

Fixes “dying ReLU” by allowing a small slope for $x < 0$.

When?

Same places as ReLU, but safer if you notice many dead neurons.

ReLU (Rectified Linear Unit)



Piece-wise Linear


Advantages of both linear and non-linear property

$$f(x) = \max(0, x)$$

Overcome the
Vanishing Gradient
Problem


$$\frac{\partial f(x)}{\partial x} = \begin{cases} 1, & x > 0 \\ 0, & x < 0 \end{cases}$$

Variations of ReLU



Leaky ReLU

$$f(x) = \max(0.01 * x, x)$$



ELU (Exponential Linear Unit)

$$f(x) = \max(\alpha * (e^x - 1), x)$$

Tanh Function

Formula:

$$f(x) = \tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

Derivative:

$$f'(x) = 1 - \tanh^2(x)$$

Range: $(-1, 1)$

Intuition: Similar to sigmoid, but centered at zero (better for optimization).

Problem: Still suffers from vanishing gradients.

Use:

Outputs in $[-1,1]$ \rightarrow zero-centered, often better than sigmoid.

Historically popular in RNNs.

Limitations:


Still suffers from vanishing gradients.

When?

Good choice in hidden layers **if you need negative outputs.**

Sometimes in old RNN architectures (before LSTMs/GRUs).

Tanh Function



$$\tanh = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

Sigmoid Function

Formula:

$$f(x) = \frac{1}{1 + e^{-x}}$$

Derivative:

$$f'(x) = f(x)(1 - f(x))$$

Range: (0,1)

Intuition: Smooth "S" curve → squashes values to probability-like output.

Use Case: Logistic regression, binary classification.

Problems: Vanishing gradients for very large/small x, slow convergence.

Use:

Outputs between **0 and 1** → good for **probabilities**.

Still used in **binary classification (output layer)**.

Limitations:


Causes **vanishing gradients** for large $|x|$.

When?

Final output layer of binary classification models.


Logistic regression.


Sigmoid Function



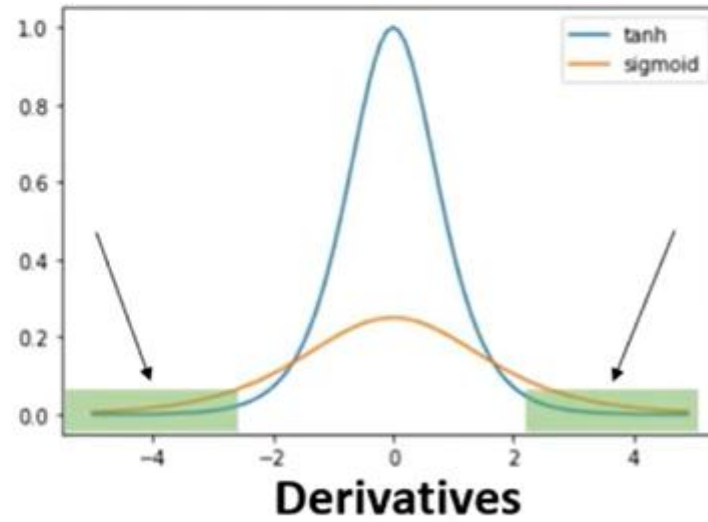


Sigmoid

$$f(x) = \frac{1}{1 + e^{-x}}$$






Derivatives




Smaller value of
derivative leads to
very slow learning

Vanishing Gradient Problem



Binary Classification



Multi-Class Classification

??

Softmax (for output layers in classification)

Formula:

$$f(x_i) = \frac{e^{x_i}}{\sum_j e^{x_j}}$$

Derivative (for backprop):

$$\frac{\partial f(x_i)}{\partial x_j} = f(x_i)(\delta_{ij} - f(x_j))$$

Intuition: Converts raw scores into probabilities across classes.


Use:

Converts a vector of logits into a **probability distribution**.

Used only in **output layer of multi-class classification**.

When?

Final layer for multi-class problems (e.g., image classification with 10 labels).



$$\begin{bmatrix} 7 \\ 0.5 \\ 1 \\ 3 \end{bmatrix}$$




$$\begin{bmatrix} e^7 \\ e^{0.5} \\ e^1 \\ e^3 \end{bmatrix}$$



$$\begin{bmatrix} \frac{e^7}{e^7 + e^{0.5} + e^1 + e^3} \\ \frac{e^{0.5}}{e^7 + e^{0.5} + e^1 + e^3} \\ \frac{e^1}{e^7 + e^{0.5} + e^1 + e^3} \\ \frac{e^3}{e^7 + e^{0.5} + e^1 + e^3} \end{bmatrix}$$



$$\begin{bmatrix} 0.9782 \\ 0.00147 \\ 0.00242 \\ 0.0179 \end{bmatrix}$$



Why *mostly* these few activation functions?

- **History & Experiments:**

Many activations (e.g., Gaussian, sine, polynomial) were tried in early neural networks. But they caused:

- Very slow training.
 - Vanishing or exploding gradients.
 - Difficulty in generalizing.
- Through trial-and-error, researchers found **Sigmoid, Tanh, ReLU, and their variants** worked best in practice.

- **Practical Needs:**

- We need functions that are **non-linear**, otherwise the whole network is just a linear regression.
- They must be **differentiable** (or piecewise differentiable) so gradient descent can work.
- Their **derivatives** should not vanish or explode for most input ranges.

Activation Functions Demo with Formulas & Explanations

```
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
```

Range of x values

```
x = np.linspace(-10, 10, 400)
```

Functions

```
step = np.where(x >= 0, 1, 0)
```

```
sigmoid = 1 / (1 + np.exp(-x))
```

```
tanh = np.tanh(x)
```

```
relu = np.maximum(0, x)
```

```
leaky_relu = np.where(x > 0, x, 0.1*x)
```

Step Function

Sigmoid

Tanh

ReLU


Leaky ReLU

Softmax needs 2D input → shape (1, n)


```
softmax = tf.nn.softmax(x.reshape(1, -1)).numpy().flatten()
```

```
# Plotting
plt.figure(figsize=(15,10))


# Step
plt.subplot(2,3,1)
plt.plot(x, step, label=r'Step:  $f(x)=1$  if  $x \geq 0$ ,  $0$  otherwise')
plt.title("Step Function")
plt.grid(True); plt.legend()
```



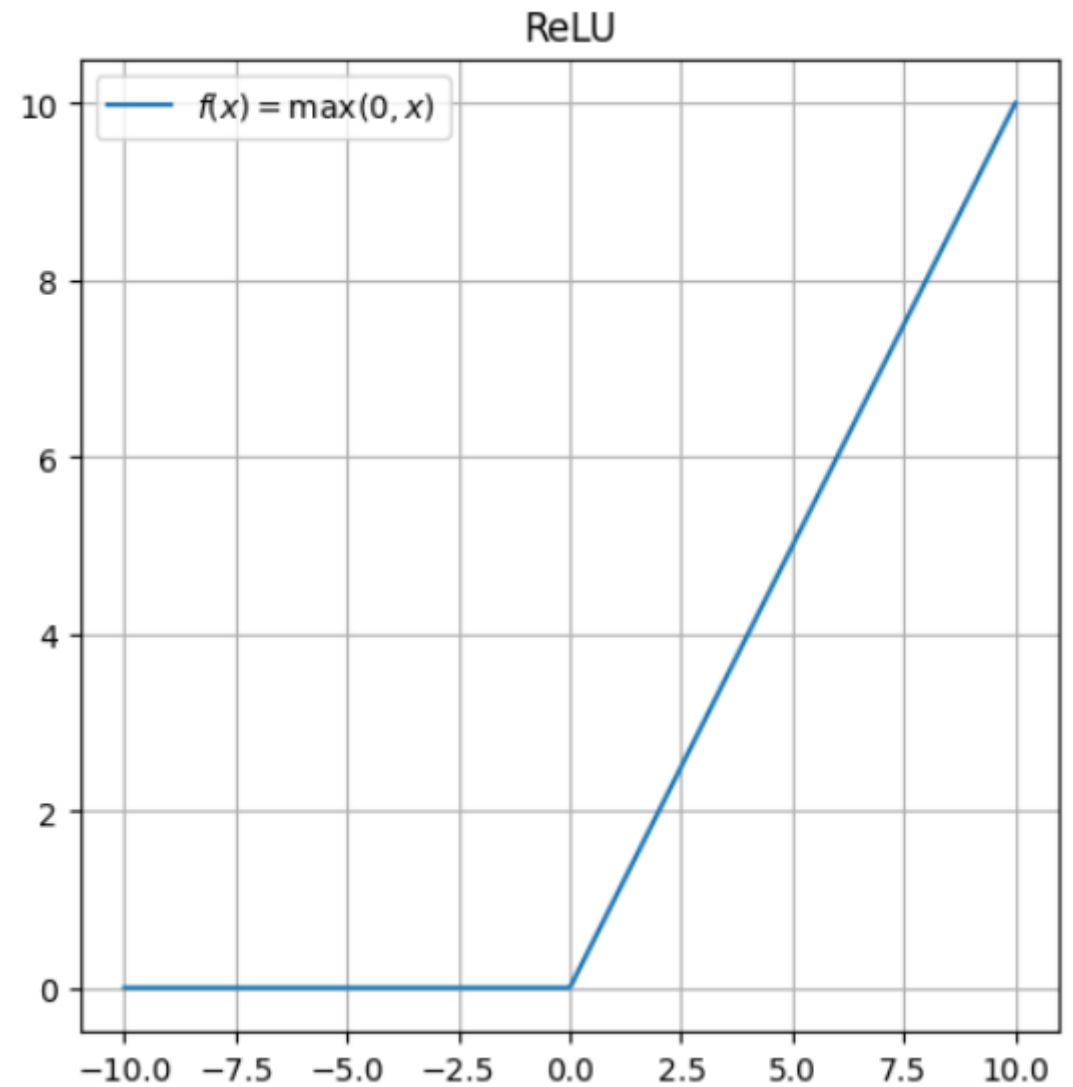
```
# Sigmoid
plt.subplot(2,3,2)
plt.plot(x, sigmoid, label=r'$\sigma(x)=\frac{1}{1+e^{-x}}$')
plt.title("Sigmoid")
plt.grid(True); plt.legend()
```




```
# Tanh
plt.subplot(2,3,3)
plt.plot(x, tanh, label=r'$\tanh(x)=\frac{e^x - e^{-x}}{e^x + e^{-x}}$')
plt.title("Tanh")
plt.grid(True); plt.legend()
```



```
# ReLU
plt.subplot(2,3,4)
plt.plot(x, relu, label=r'$f(x)=\max(0,x)$')
plt.title("ReLU")
plt.grid(True); plt.legend()
```




```
# Leaky ReLU
plt.subplot(2,3,5)
plt.plot(x, leaky_relu, label=r'Leaky ReLU:  $f(x)=x$  if  $x>0$ , else  $0.1x$ ')
plt.title("Leaky ReLU")
plt.grid(True); plt.legend()
```




```
# Softmax
plt.subplot(2,3,6)
plt.plot(x, softmax, label=r'$f(x_i)=\frac{e^{x_i}}{\sum_j e^{x_j}}$')
plt.title("Softmax")
plt.grid(True); plt.legend()




plt.tight_layout()
plt.show()
```



1. **Step Function**: Outputs 0 or 1. Used in early perceptrons but not differentiable.
2. **Sigmoid**: Smoothly maps input to (0,1). Good for probabilities but causes vanishing gradients.
3. **Tanh**: Maps input to (-1,1). Zero-centered but still suffers vanishing gradients.
4. **ReLU**: Outputs positive values as is, else 0. Very popular, avoids vanishing gradients (mostly).
5. **Leaky ReLU**: Like ReLU but allows small negative slope. Solves 'dying ReLU' problem.
6. **Softmax**: Converts vector into probability distribution. Common in output layers for classification.


Rule of Thumb:

- **Hidden layers** → ReLU (or Leaky ReLU if dead neurons problem).
- **Output layer** →
 - **Sigmoid** for binary classification.
 - **Softmax** for multi-class classification.
 - **No activation (linear)** for regression.



2) Either ReLU or tanh can be used in hidden layers

1) Sigmoid in output Neuron for Binary Classification



Activation	Formula	Range	Pros	Cons	When to Use
Step	$f(x) = 1$ if $x \geq 0$, else 0	{0,1}	Simple, mimics biological neurons	Not differentiable → can't train with backprop	Rarely; only toy models or binary hard thresholds
Sigmoid	$\sigma(x) = \frac{1}{1+e^{-x}}$	(0,1)	Smooth, interpretable as probability	Vanishing gradients, not zero-centered	Output layer of binary classification
Tanh	$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$	(-1,1)	Zero-centered, stronger gradients than sigmoid	Still vanishing gradient	Sometimes in hidden layers ; classic RNNs
ReLU	$f(x) = \max(0, x)$	$[0, \infty)$	Simple, efficient, avoids vanishing gradients	"Dying ReLU" (neurons stuck at 0)	Default for hidden layers in deep nets
Leaky ReLU	$f(x) = x$ if $x > 0$, else $0.01x$	$(-\infty, \infty)$	Fixes dying ReLU, allows small gradient when $x < 0$	Small slope may affect learning	Hidden layers if ReLU has many dead neurons
Softmax	$f(x_i) = \frac{e^{x_i}}{\sum_j e^{x_j}}$	(0,1), sums to 1	Produces probability distribution	Sensitive to large input values	Output layer for multi-class classification
Linear (Identity)	$f(x) = x$	$(-\infty, \infty)$	Keeps values unchanged	No non-linearity	Regression tasks (output layer)

Role of Derivatives

The **derivative of the activation** is what flows backward during backpropagation:

$$\frac{\partial L}{\partial z} = \frac{\partial L}{\partial a} \cdot f'(z)$$

If $f'(z)$ is too small (like sigmoid near saturation), the gradient vanishes → **no learning**.

If $f'(z)$ is huge, gradient explodes → **unstable training**.

```
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf

# Input range
x = np.linspace(-10, 10, 500, dtype=np.float32)

# Activation functions
sigmoid = tf.nn.sigmoid(x).numpy()
tanh = tf.nn.tanh(x).numpy()
relu = tf.nn.relu(x).numpy()
leaky_relu = tf.nn.leaky_relu(x, alpha=0.1).numpy()

# Derivatives
# Sigmoid derivative:  $\sigma(x) * (1 - \sigma(x))$ 
sigmoid_deriv = sigmoid * (1 - sigmoid)

# Tanh derivative:  $1 - \tanh(x)^2$ 
tanh_deriv = 1 - np.power(tanh, 2)



# ReLU derivative: 0 if  $x < 0$  else 1
relu_deriv = np.where(x > 0, 1, 0)

# Leaky ReLU derivative: alpha if  $x < 0$  else 1
leaky_relu_deriv = np.where(x > 0, 1, 0.1)
```

```
# Plot
plt.figure(figsize=(14,10))

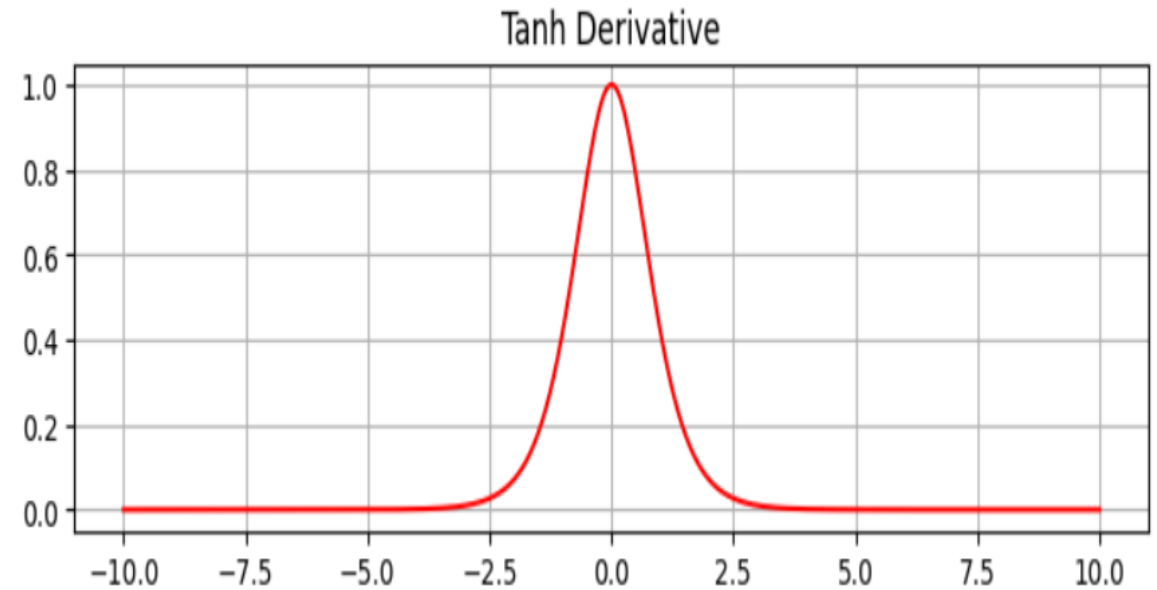

# Sigmoid
plt.subplot(4,2,1)
plt.plot(x, sigmoid, label="Sigmoid")
plt.title("Sigmoid Activation")
plt.grid(True)

plt.subplot(4,2,2)
plt.plot(x, sigmoid_deriv, label="Sigmoid'", color="red")
plt.title("Sigmoid Derivative")
plt.grid(True)
```




```
# Tanh
plt.subplot(4,2,3)
plt.plot(x, tanh, label="Tanh")
plt.title("Tanh Activation")
plt.grid(True)


plt.subplot(4,2,4)
plt.plot(x, tanh_deriv, label="Tanh'", color="red")
plt.title("Tanh Derivative")
plt.grid(True)
```




```
# ReLU
plt.subplot(4,2,5)
plt.plot(x, relu, label="ReLU")
plt.title("ReLU Activation")
plt.grid(True)

plt.subplot(4,2,6)
plt.plot(x, relu_deriv, label="ReLU'", color="red")
plt.title("ReLU Derivative")
plt.grid(True)
```

ReLU Activation





ReLU Derivative



```
# Leaky ReLU
plt.subplot(4,2,7)
plt.plot(x, leaky_relu, label="Leaky ReLU")
plt.title("Leaky ReLU Activation")
plt.grid(True)

plt.subplot(4,2,8)
plt.plot(x, leaky_relu_deriv, label="Leaky ReLU'", color="red")
plt.title("Leaky ReLU Derivative")
plt.grid(True)

plt.tight_layout()
plt.show()
```



Sigmoid → outputs between $(0,1)$, but derivative becomes very small for $|x| > 5$ → vanishing gradient problem.

Tanh → centered around 0, derivative stronger near 0, but still vanishes for large $|x|$.

ReLU → derivative is 0 for negative x (dead neurons) but 1 for positive x → efficient gradient flow.

Leaky ReLU → fixes dead ReLU by keeping a small slope on negative side.

THANK YOU